

# Traffic Sign Classifier

Jianxing Ke

December 27, 2018

The goals / steps of this project are the following:

- Load the data set (see below for links to the project data set)
- Explore, summarize and visualize the data set
- Design, train and test a model architecture
- Use the model to make predictions on new images
- Analyze the softmax probabilities of the new images
- Summarize the results with a written report

## 1 Data Set Summary & Exploration

### 1.1 Summary of the data set

The data set used in this project was *German Traffic Sign Dataset*. After the data was loaded, we see that the size of training examples is 34799; the size of testing examples is 12630; the size of validation examples is 4410. The dataset contains 43 different traffic signs, and each image is 32x32x3.

To visualize the dataset, I randomly output 10 images, and counted the number of examples of each class and displayed it as a bar graph.

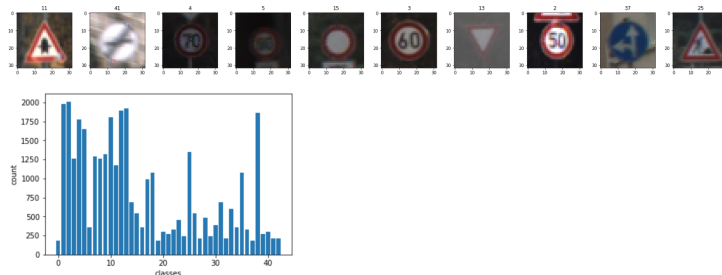


Figure 1: Image examples and bar graph

## 2 Model Architecture

### 2.1 Data Preprocessing

With the given dataset, I generated augmented data in the following ways:  
**Scale images.** For each image, I scaled it with factor 0.75 and 0.6 with respect to the center of the image, then resized it to 32x32x3.




Figure 2: Scaled images 1




Figure 3: Scaled images 2

**Rotate images.** For each image, I rotated it 15 degrees clockwise and counterclockwise, without changing the size of the image. Here are two example outputs.




Figure 4: Rotated images 1




Figure 5: Rotated images 2

**Add a value to images.** For each image, I added a value of 3 and -3 to each pixels. The output images look similar to the original.




Figure 6: Images with added value

**Normalize images.** As suggested in the project instructions, for each image, I normalized it by subtracting 128 and divided by 128. By doing this the dataset would have 0 mean and equal variance.




Figure 7: Normalized images

**Turn images to grayscale.** For our task of classifying the traffic signs, color is not really a useful information since we focus more on different shapes on a sign.




Figure 8: Grayscaled images

After all the preprocessing steps, I had a dataset of 243593 images, each has dimension  $32 \times 32 \times 1$ .

## 2.2 Model Architecture

My current model is very similar to the LeNet model. I think my model underfitted for this project, even though it achieved pretty decent accuracy with the augmented dataset.

The model had the following layers:

Layer	Description
Input	$32 \times 32 \times 1$ gray image
Convolution	16 $5 \times 5$ filters, stride 1x1, valid padding, outputs $28 \times 28 \times 16$
RELU	
Max pooling	$2 \times 2$ filter, stride $2 \times 2$ , valid padding, outputs $14 \times 14 \times 16$
Convolution	16 $5 \times 5$ filters, stride 1x1, valid padding, outputs $10 \times 10 \times 16$
RELU	
Max pooling	$2 \times 2$ filter, stride $2 \times 2$ , valid padding, outputs $5 \times 5 \times 16$
Fully Connected	input $5 \times 5 \times 16$ , outputs 120
RELU	
Fully Connected	input 120, outputs 84
RELU	
Output	input 84, outputs 43

## 2.3 Model Training

To train the model, I used Adam optimizer with learning rate 0.001, and the loss function was cross entropy. During each epoch, the dataset was randomly shuffled, then a batch was fed into the classifier at a time. With the number of epochs 10 and batch size 256, my model had an validation accuracy of 94.9%. Without the augmented data, the model only achieved about 90% validation accuracy. Compared to other well-known architectures that achieve 99% accuracy, more layers can be added to my model, and parameters can be better tuned to achieve better results.

### 3 Test on New Images

To test my model, I picked 5 extra traffic signs online.



Figure 9: test images

These images were in different sizes, so I needed to preprocess them before feeding them to the model. Below are images that were actually fed to the model.




Figure 10: Preprocessed test images

The model achieved 80% accuracy on the 5 test images. To find out which image the model predicted wrong, I calculated the top 5 softmax probabilities for test images.



Figure 11: Top 5 softmax probabilities

From this, we see that the model predicted 4 traffic signs correctly, with high confidence. But it also predicted 100km/h speed limit sign incorrectly with high confidence.

## 4 Visualize the Network's State

With the given `outputFeatureMap` function, we can easily visualize a model's feature layers. I plot the outputs of my model's 2 convolution layers for one of the test images here.




Figure 12: Conv layer 1




Figure 13: Conv layer 2

The feature maps captured the shape (triangle) of the sign very well, then it also captured the contents of the sign, in this case two human figures.